



Private offline AI assistant for financial services

Select a model ?

mistral

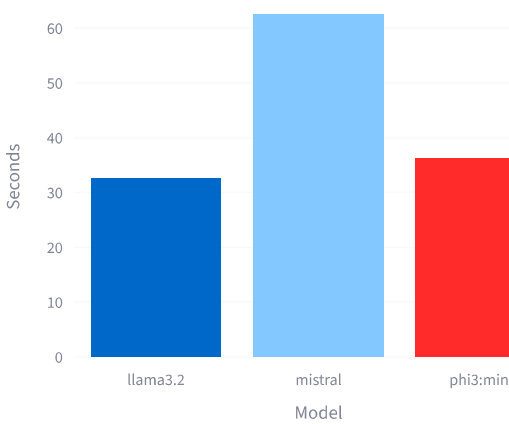
Local model — no data leaves your device

Built with Ollama + FastAPI + Streamlit

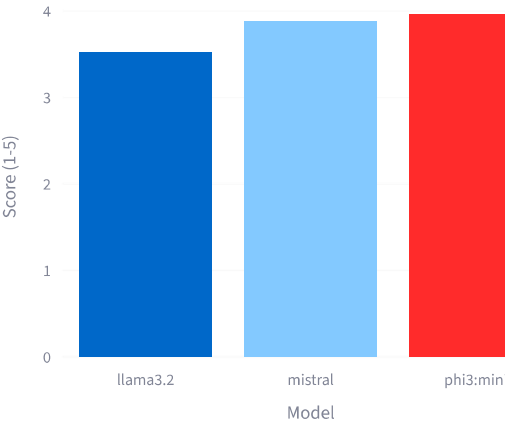
Benchmark Results

Run benchmark.py first to see results here

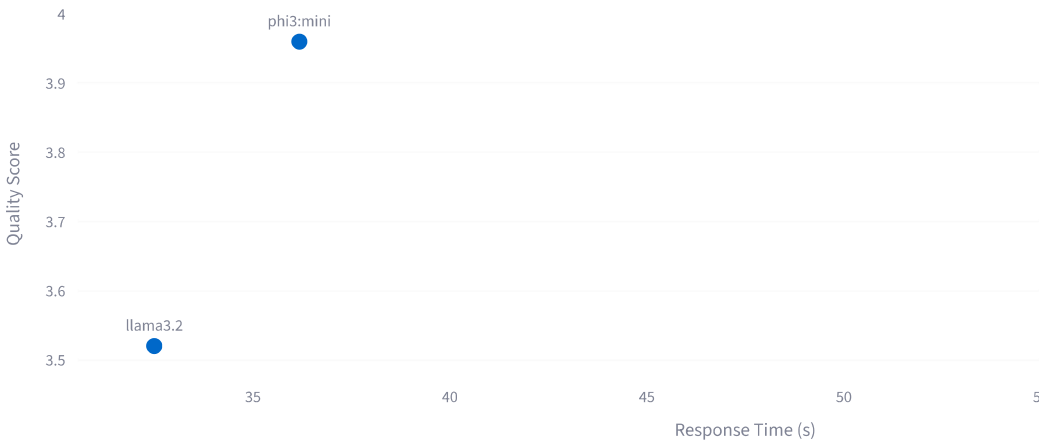
Average Response Time (lower is better)



Average Quality Score (higher is better)



Speed vs Quality Tradeoff



Results by Category

Quality Score by Category and Model

